

Application Layer

Dr. A Krishna Chaitanya,
Indian Institute of Information Technology Sri City

30-Aug-2017

- Applications use the services of network (Transport layer)
- For an application developer, architecture and services of network are fixed
- Architectures of applications:
 - **Client-Server** architecture
 - **Peer-to-Peer** (P2P) architecture
- Application developer decides on the architecture and services of transport layer to be used.

Client-Server Architecture

- Server: An end system that **serves the requests** from various hosts.
- A server is always **ON**.
- Client: An end system that **requests** a server for content.
- A client can be either **ON-OFF** or always **ON**.
- Example applications using this architecture: web, e-mail, file transfer, etc.

Peer-to-Peer Architecture

- End systems communicate by a direct connection.
- The end systems are called peers.
- Example applications: skype, internet telephony, torrents, etc
- Advantages:
 - File distribution
 - Self-scalable: can handle growth in traffic
 - Cost effective: no server infrastructure and server bandwidth.
- Challenges in P2P Architecture:
 - ISP friendly: asymmetric data traffic.
 - Security
 - Incentives: Peers should share bandwidth.

Processes Communicating

- A process is a program that is running within an end system.
- A client process is a process running on a client and a server process is process running on a server.
- It is the client process and server processes that are actually communicating.
- A process sends and receives messages to and from transport layer through a software interface known as **socket**.
- A socket is also known as **Application Programming Interface (API)**.

Services of Transport Layer

- **Reliable data transfer**: Guaranteed data delivery service.
- **Throughput**
- **Timing**: for example, it is guaranteed that a packet will be delivered no more than 100 msec later.
- **security**: end-point authentication, encryption and decryption.

- Transmission Control Protocol (TCP)
 - Connection oriented service: handshaking, full-duplex connection
 - Reliable data transfer service: packets get delivered without error and in proper order.
 - Congestion control
- User Datagram Protocol (UDP)
 - Connectionless
 - Unreliable data transfer service.
 - No congestion control

Applications

Application	Application-Layer Protocol	Underlying Transport Protocol
Electronic mail	SMTP [RFC 5321]	TCP
Remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
File transfer	FTP [RFC 959]	TCP
Streaming multimedia	HTTP (e.g., YouTube)	TCP
Internet telephony	SIP [RFC 3261], RTP [RFC 3550], or proprietary (e.g., Skype)	UDP or TCP

Addressing Processes

- There are many processes running on a host, how to identify the destination process?

Addressing Processes

- There are many processes running on a host, how to identify the destination process?
- We identify host by IP address.

Addressing Processes

- There are many processes running on a host, how to identify the destination process?
- We identify host by **IP address**.
- We identify processes by **port numbers**!
- For example, web server is identified by port number 80, mail server is identified by port number 25.

- A web page is a document and consists of objects
- An object is nothing but a file such as HyperText Markup Language (HTML) file, an image file, applet or video clip.
- If a web page contains a basic html file and ten images, we say the web page contains 11 objects.
- HyperText Transfer Protocol (HTTP) is the web's application layer protocol
- HTTP uses client-server architecture with TCP.
- The client program and server program talk to each other by exchanging HTTP messages.

Uniform Resource Locator

- An object should be addressable by a URL.
- Each URL consists of hostname and objects path name
- For example, `http://www.iiits.ac.in/wp-content/uploads/2017/05/Untitled-design-15.png` is url for an image.
- `www.iiits.ac.in` is host name
- `wp-content/uploads/2017/05/Untitled-design-15.png` is path name.
- Client side of HTTP is implemented in Web browser and server side is implemented in Web server.
- Examples: Apache and Microsoft Internet Information server.

- HTTP client initiates a connection with HTTP server (**handshaking**).
- Once the connection is established, client and server exchange messages through socket interface.
- Client sends an HTTP request and receives HTTP messages through its socket
- Server receives HTTP requests and sends HTTP responses through its socket interface.
- Client/server need not worry about packets (does not have any control) after sending through their socket.
- Server sends requested files without storing state information of client. Thus HTTP is a **stateless** protocol.

HTTP Connection

- Let us say, a web page has one html file and 10 images.
- How does client retrieve the web page?

HTTP Connection

- Let us say, a web page has one html file and 10 images.
- How does client retrieve the web page?
- **Nonpersistent** and **Persistent**

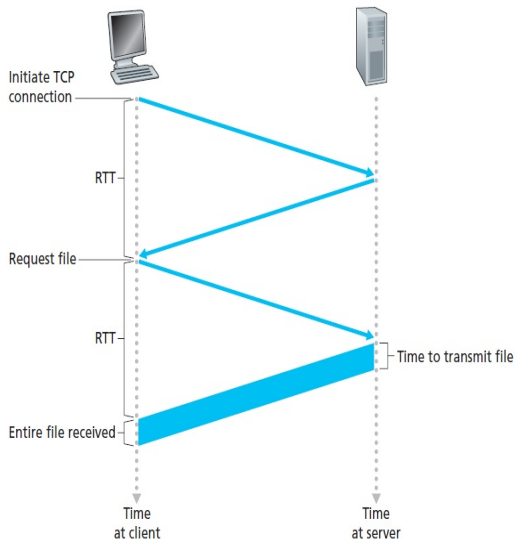
HTTP Connection

- Let us say, a web page has one html file and 10 images.
- How does client retrieve the web page?
- **Nonpersistent** and **Persistent**
- Nonpersistent: one TCP connection for **each** file
- Persistent: one TCP connection for **all** files

Nonpersistent Connection

- For each file:
 - HTTP client initiates a TCP connection to the server on port number 80
 - Client sends its HTTP request and it includes the path name to the file
 - HTTP server receives the request and retrieves the file and sends the HTTP response to the client
 - HTTP server tells TCP to close the connection.
- TCP connections can be **serial or parallel** depending on browser's configuration

Round-Trip Time



Persistent Connection

- Server leaves the connection after sending the HTTP response
- **Pipelining**: A browser can request for files without waiting for the reception of pending requests.
- TCP closes after some idle period
- Default mode HTTP: Persistent connection with pipelining.

HTTP Request Format

- HTTP request message:

GET /somedir/page.html HTTP/1.1

Host: www.iitm.ac.in

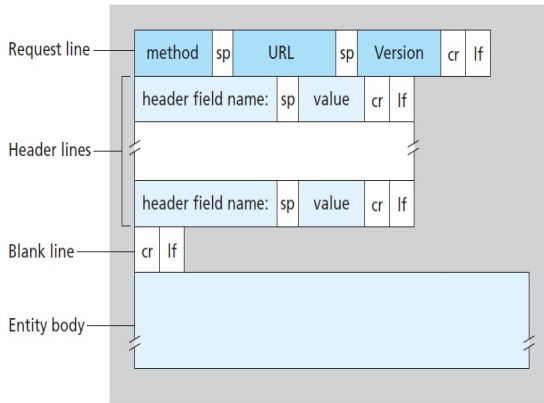
Connection: close

User-agent: Mozilla/4.0

Accept-language: En

- Methods: GET, PUT, POST, HEAD, DELETE

HTTP Request



HTTP Response

- HTTP response message:

HTTP/1.1 200 OK

Connection: close

Date: Sat, 07 Jul 2007 12:00:15 GMT

Server: Apache/1.3.0 (Unix)

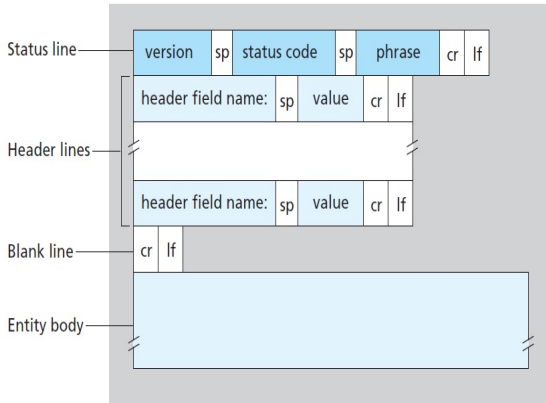
Last-Modified: Sun, 6 May 2007 09:23:24 GMT

Content-length: 6821

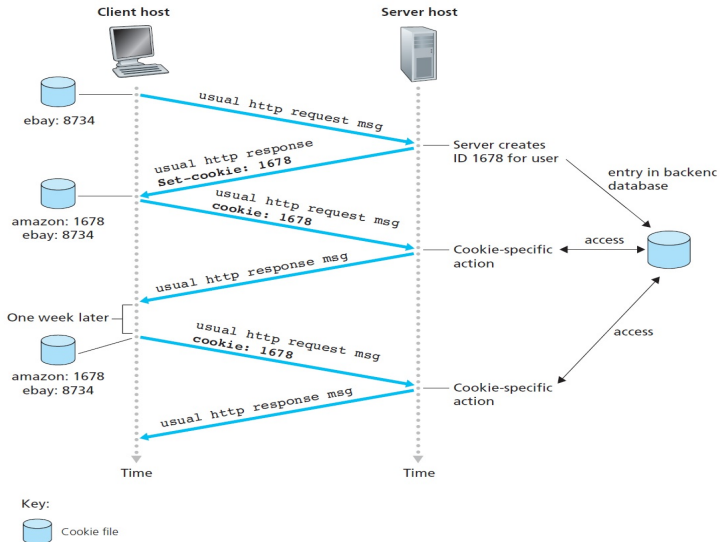
Content-Type: text/html
(data data ... data)

- 200 OK
- 301 Moved Permanently
- 404 Not Found

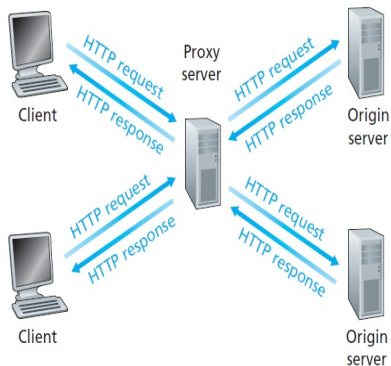
HTTP Response



Cookies



Web Caching



- Cache request message:

GET /somedir/student.jpg HTTP/1.1

Host: www.iiits.ac.in

If-modified-since: Wed, 23 Aug 2017 17:30:00

Conditional GET

- Cache request message:

GET /somedir/student.jpg HTTP/1.1

Host: www.iiits.ac.in

If-modified-since: Wed, 23 Aug 2017 17:30:00

- Response: HTTP/1.1 304 Not Modified

Date: Tue, 29 Aug 2017 13:00:00

Server: Apache/1.3.0 (Unix)

(empty empty empty)

File Transfer Protocol

- Similar to HTTP: client-server architecture, transmission control protocol
- Two parallel TCP connections to transfer a file: **TCP control connection** and **TCP data connection**
- Control information:
 - User identification
 - Change remote directory
 - Commands to **put** and **get** files
- FTP is said to control information **out-of-band** where as HTTP is said to control information **in-band**.

- Commands:
 - **USER** username
 - **PASS** password
 - **LIST**
 - **RETR** filename
 - **STOR** filename
- Replies:
 - **331** username OK, password required
 - **125** data connection already open; transfer starting
 - **425** can not open data connection
 - **452** error writing file

- Asynchronous communication medium
- Major components of e-mail system:
 - **User agent**: allows users to read, forward, save and compose messages
 - **Mail server**
 - **SMTP**
- Examples of user agents: Microsoft Outlook, Mozilla Thunderbird, Apple Mail

- User agent sends message to user's mail server.
- SMTP transfers message from user's mail server to recipient's mail server.
- Client side of SMTP is running on sender's mail server and server side of SMTP is running on recipient's mail server.
- Recipient's mail server delivers the message in recipient's mail box.

SMTP Sequence of Operations

- Alice composes message using her user agent. Provides Bob's mail address and instructs to send the message.
- User agent sends the message to her mail server and message waits in the queue of the server.
- SMTP client sees the message in the mail server and it opens a TCP connection to an SMTP server running on Bob's mail server.
- SMTP transfers the message from client to server.
- SMTP server receives the message. Bob's mail server places the message in Bob's mail box.
- Bob invokes his user agent to read the message.

SMTP Sequence of Operations

- If recipient's mail server is down, SMTP client **re-attempts** to send the message (say for every 30 minutes)
- If the delivery is not successful after some duration, it will be notified to the sender and message will be dropped.

Client-Server Conversation

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr ... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

Message Formats

- Header lines similar to those in HTTP messages
- Header must have **From:**, **To:**
- Optional header lines include **Subject:**

Comparison with HTTP

- HTTP is a **pull protocol**
- SMTP is **push protocol**
- SMTP requires each message to be 7-bit ASCII format.
HTTP does not have this restriction
- HTTP encapsulates each object in its own HTTP response message. Internet mail places all of its objects into one message.

Mail Access Protocols

- In early days of internet, Bob reads mail by logging onto mail server and executing a mail reader on that host
- Client-server architecture
- Reads e-mail by running a client on the user's end system
- Mail access protocol transfers message from Bob's mail server to his local PC.
- Popular mail access protocols: Post Office Protocol - version 3 (**POP3**), Internet Mail Access Protocol (**IMAP**) and HTTP

- Begins when a user agent opens a TCP connection with mail server on port 110.
- POP3 progresses in three phases:
 - Authorization
 - Transaction
 - Update
- Authorization: **user** <username> and **pass** <password>
- Transaction: user agent sends commands and server responds with **+OK** and **-ERR**

POP3 Transaction


- Two modes:
 - download and delete
 - download and keep
- Download and delete:

```
C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: (blah blah ...
S: .....
S: .....blah)
S: .
C: dele 1
C: retr 2
S: (blah blah ...
S: .....
S: .....blah)
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```


IMAP and HTTP

- IMAP associates each message with a folder
- Provides commands to allow users to **create folder** and **move messages across folders**
- Provides commands to search for a message
- Maintains user **state information** across IMAP sessions
- Components of messages can be retrieved
- HTTP:
 - e-mail access through web browser
 - web browser communicates to the mail server via HTTP

Problem



- Average object size is 1Mbits
- Average request rate: 15 objects per sec.
- Time from request to receive response is 2 sec on an average.

Total Response Time

Total Response Time

- Traffic intensity on the LAN

Total Response Time

- Traffic intensity on the LAN
- 0.15

Total Response Time

- Traffic intensity on the LAN
- 0.15
- Traffic intensity on the access link

Total Response Time

- Traffic intensity on the LAN
- 0.15
- Traffic intensity on the access link
- 1

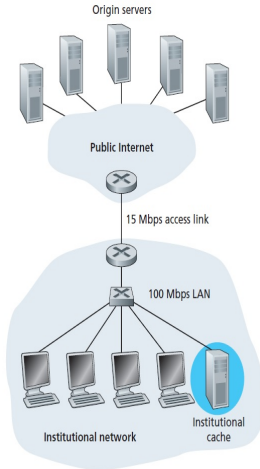
Total Response Time

- Traffic intensity on the LAN
- 0.15
- Traffic intensity on the access link
- 1
- Suppose, access link upgraded to 100Mbps link.

Total Response Time

- Traffic intensity on the LAN
- 0.15
- Traffic intensity on the access link
- 1
- Suppose, access link upgraded to 100Mbps link.
- Expensive Solution!

Problem



- Assume cache provides hit rate of 0.4

Total Response Time

- Traffic intensity on the LAN

Total Response Time

- Traffic intensity on the LAN
- 0.15

Total Response Time

- Traffic intensity on the LAN
- 0.15
- Traffic intensity on the access link

Total Response Time

- Traffic intensity on the LAN
- 0.15
- Traffic intensity on the access link
- 0.6

Total Response Time

- Traffic intensity on the LAN
- 0.15
- Traffic intensity on the access link
- 0.6
- Typically delay ten milliseconds

Total Response Time

- Traffic intensity on the LAN
- 0.15
- Traffic intensity on the access link
- 0.6
- Typically delay ten milliseconds
- Average response time: $0.4*(0.01)+0.6*(2.01)$ seconds


Applications of Peer-to-Peer Architecture

- **File distribution**: application that transfers a file from a single source to multiple peers.
- **Database distributed** over a large community of peers.
- **Internet telephony** : Skype.

File Distribution

- Each peer can **redistribute** any portion of the file to any other peer
- Popular file distribution protocol : BitTorrent, developed by Bram Cohen
- Scalability

Scalability



- N peers
- **Distribution time**: the time required to distribute a file to all peers.

Assumptions

- Internet has abundant bandwidth and all bottlenecks are in the network access
- All the server and client bandwidth is available for file distribution

Distribution Time for Client-Server Architecture

- Let D_{cs} denote the distribution time for client-server architecture for a file size of F bits
- The server has to transmit a total of NF bits at an upload rate of $u_s bps$.

Distribution Time for Client-Server Architecture

- Let D_{cs} denote the distribution time for client-server architecture for a file size of F bits
- The server has to transmit a total of NF bits at an upload rate of $u_s bps$.
- Minimum time required for distribution is $\frac{NF}{u_s}$ seconds

Distribution Time for Client-Server Architecture

- Let D_{cs} denote the distribution time for client-server architecture for a file size of F bits
- The server has to transmit a total of NF bits at an upload rate of $u_s bps$.
- Minimum time required for distribution is $\frac{NF}{u_s}$ seconds
- Let $d_{min} = \min\{d_1, \dots, d_N\}$

Distribution Time for Client-Server Architecture

- Let D_{cs} denote the distribution time for client-server architecture for a file size of F bits
- The server has to transmit a total of NF bits at an upload rate of $u_s bps$.
- Minimum time required for distribution is $\frac{NF}{u_s}$ seconds
- Let $d_{min} = \min\{d_1, \dots, d_N\}$
- Minimum distribution time is $\frac{F}{d_{min}}$ seconds

Distribution Time for Client-Server Architecture

- Let D_{cs} denote the distribution time for client-server architecture for a file size of F bits
- The server has to transmit a total of NF bits at an upload rate of $u_s bps$.
- Minimum time required for distribution is $\frac{NF}{u_s}$ seconds
- Let $d_{min} = \min\{d_1, \dots, d_N\}$
- Minimum distribution time is $\frac{F}{d_{min}}$ seconds
- Thus,

$$D_{cs} \geq \max \left\{ \frac{NF}{u_s}, \frac{F}{d_{min}} \right\}$$

Distribution Time for Client-Server Architecture

- Let D_{cs} denote the distribution time for client-server architecture for a file size of F bits
- The server has to transmit a total of NF bits at an upload rate of u_s bps.
- Minimum time required for distribution is $\frac{NF}{u_s}$ seconds
- Let $d_{min} = \min\{d_1, \dots, d_N\}$
- Minimum distribution time is $\frac{F}{d_{min}}$ seconds
- Thus,

$$D_{cs} \geq \max \left\{ \frac{NF}{u_s}, \frac{F}{d_{min}} \right\}$$

- Show that

$$D_{cs} = \max \left\{ \frac{NF}{u_s}, \frac{F}{d_{min}} \right\}$$

Distribution Time for P2P Architecture

- The server has to send each bit of the file at least once:
Minimum distribution time is at least $\frac{F}{u_s}$ seconds

Distribution Time for P2P Architecture

- The server has to send each bit of the file at least once:
Minimum distribution time is at least $\frac{F}{u_s}$ seconds
- The peer with lowest download rate can not obtain F bits in less than $\frac{F}{d_{min}}$ seconds

Distribution Time for P2P Architecture

- The server has to send each bit of the file at least once:
Minimum distribution time is at least $\frac{F}{u_s}$ seconds
- The peer with lowest download rate can not obtain F bits in less than $\frac{F}{d_{min}}$ seconds
- The total upload rate $u_{total} = u_s + u_1 + \dots + u_N$. The system must deliver F bits to each of the N peers: Minimum distribution time is $\frac{NF}{u_{total}}$

Distribution Time for P2P Architecture

- The server has to send each bit of the file at least once:
Minimum distribution time is at least $\frac{F}{u_s}$ seconds
- The peer with lowest download rate can not obtain F bits in less than $\frac{F}{d_{min}}$ seconds
- The total upload rate $u_{total} = u_s + u_1 + \dots + u_N$. The system must deliver F bits to each of the N peers: Minimum distribution time is $\frac{NF}{u_{total}}$
- Thus, minimum distribution time D_{P2P} is at least

$$\max\left\{\frac{F}{u_s}, \frac{F}{d_{min}}, \frac{NF}{u_{total}}\right\}$$

Distribution Time for P2P Architecture

- The server has to send each bit of the file at least once:
Minimum distribution time is at least $\frac{F}{u_s}$ seconds
- The peer with lowest download rate can not obtain F bits in less than $\frac{F}{d_{min}}$ seconds
- The total upload rate $u_{total} = u_s + u_1 + \dots + u_N$. The system must deliver F bits to each of the N peers: Minimum distribution time is $\frac{NF}{u_{total}}$
- Thus, minimum distribution time D_{P2P} is at least

$$\max\left\{\frac{F}{u_s}, \frac{F}{d_{min}}, \frac{NF}{u_{total}}\right\}$$

- Assumption: each peer can redistribute a bit as soon as it receives the bit.


Distribution Time for P2P Architecture

- The server has to send each bit of the file at least once:
Minimum distribution time is at least $\frac{F}{u_s}$ seconds
- The peer with lowest download rate can not obtain F bits in less than $\frac{F}{d_{min}}$ seconds
- The total upload rate $u_{total} = u_s + u_1 + \dots + u_N$. The system must deliver F bits to each of the N peers: Minimum distribution time is $\frac{NF}{u_{total}}$
- Thus, minimum distribution time D_{P2P} is at least

$$\max\left\{\frac{F}{u_s}, \frac{F}{d_{min}}, \frac{NF}{u_{total}}\right\}$$

- Assumption: each peer can redistribute a bit as soon as it receives the bit.
- There is a scheme that actually achieves this lower bound.

Distribution Time for P2P Architecture




- All peers upload at a rate of u bps.
- $\frac{F}{u} = 1$ hour, $u_s = 10u$ and $d_{min} \geq u_s$.

- Collection of peers participating in the distribution of a file is called a **torrent**
- Peers in a torrent download equal-size **chunks** of the file (typically 256 KBytes)
- A peer accumulates more and more chunks over time
- Once a peer has acquired complete file, it may leave the torrent or continue to participate in the torrent
- Peers may leave torrents with subsets of chunks

Bit Torrent

- Each torrent has a node called **tracker**.
- When a peer joins the torrent, it registers with the tracker
- Each peer in the torrent **periodically updates the tracker** about its presence.



- Alice receives a subset of participating peers in the torrent
- She establishes TCP connection with some of the peers and we call them as **neighboring peers of Alice**
- Neighboring peers may vary over time
- Each peer will have some subset of chunks from the file, with different peers having different subsets
- Alice maintains a list of chunks that her neighbors have.

- Alice will issue requests for chunks she currently does not have
- Which chunks should be requested first?

- Alice will issue requests for chunks she currently does not have
- To which of her neighbors should she send requested chunks?

- Alice will issue requests for chunks she currently does not have
- Which chunks should be requested first?
- Rarest first: finds the chunks that are rarest among her neighbors

- Alice will issue requests for chunks she currently does not have
- To which of her neighbors should she send requested chunks?
- Tit-for-tat

- Alice gives priority to the neighbors that are currently supplying her data at the highest rate
- Every 30 seconds, she also picks one additional neighbor at random and sends it chunks. Let it be Bob.
- In due course of time, Alice, may become one of the top uploaders in which case Bob could start sending data to Alice.

Distributed Hash Tables (DHT)

- Huge database to be stored among number of peers in a distributed way
- Database is consists of (key, value) pairs. For Example, (PAN No., Aadhar No.), (Content Name, IP), etc.
- Peers query the database by supplying the key and database replies the matching pairs to the querying peer
- How to store database among the peers


- Assign an **identifier** to each peer.
- An identifier is an integer in $[0, 2^n - 1]$ for some fixed n
- (key, value) pairs are also identified by integers using **hash functions**
- Hash function is available to all peers.

- Define a rule for assigning keys to peers

- Define a rule for assigning keys to peers
- Closest to the key:
- For example, $n = 4$, with eight peers: 1,3,4,5,8,10,12 and 15.
Store (11, 0123-4567-8910) in one of the eight peers


- Define a rule for assigning keys to peers
- Closest to the key:
- For example, $n = 4$, with eight peers: 1,3,4,5,8,10,12 and 15. Store (11, 0123-4567-8910) in one of the eight peers
- By closest convention, peer 12 is the immediate successor for key 11. Store in peer 12.
- If the key is larger than all the peer identifiers, we use modulo- 2^n convention.

Circular DHT



- Each peer is aware of only its immediate predecessor and successor
- N messages at most

Shortcut



- Number of shortcuts are relatively small in number
- How many shortcut neighbors and which peers should be these shortcut neighbors? Research problem: $O(\log(N))$

Peer Churn

- Peers can come and go without warning
- Peers keep track to two immediate predecessor and successors.
- When a peer abruptly leaves, its predecessor and successor learn that a peer has left and **updates the list of its predecessor and successor**.

